

# Confluent Certified Developer for Apache Kafka Certification Examination -Notes

<b>Setting Up Environment</b>	<b>2</b>
Running Apache Kafka in docker container - First time	3
Stopping Apache Kafka that is running in docker container	4
Starting the container that was stopped	5
<b>Deleting Kafka environment</b>	<b>5</b>
Troubleshooting	5
Checking state of kafka process that are running in docker container	5
Stop the Linux virtual machine	6
<b>Application Design</b>	<b>6</b>
Kafka's command line tools	6
Topic Operations	6
Creating a New Topic in a cluster	6
Key Points	7
Increasing partitions of topic	7
Key Points	7
Decreasing partitions of topic	7
Deleting a Topic	7
Key point	8
Listing all topics in a Cluster	8
Describe all the topics	8
Console Producer	8
Console Consumer	8
Consumer Groups	9
List Consumer Groups in a cluster	9
Describe a Consumer Group	9
<b>Development</b>	<b>9</b>
<b>Deployment/Testing/Monitoring</b>	<b>9</b>
<b>References</b>	<b>9</b>

# Setting Up Environment

We will be using confluent kafka docker images to run the kafka services. Follow

## Running Apache Kafka in docker container - First time

- Clone the repository

- 

```
# git clone  
https://github.com/balajich/CCD-Apache-Kafka-Certification-Examination-Notes.git
```

- Change to cloned directory

- 

```
# cd CCD-Apache-Kafka-Certification-Examination-Notes
```

- Start the linux virtual machine using Vagrant

- 

```
# vagrant up
```

- Take ssh to linux

- 

```
# vagrant ssh
```

- Switch to root user in linux machine

- 

```
[vagrant@kserver ~]$ sudo su -
```

- Go to vagrant folder

- 

```
[root@kserver ~]# cd /vagrant/
```

- Start the kafka server using docker-compose, The below command will start all the dependent services

- 

```
[root@kserver vagrant]# docker-compose up -d
```

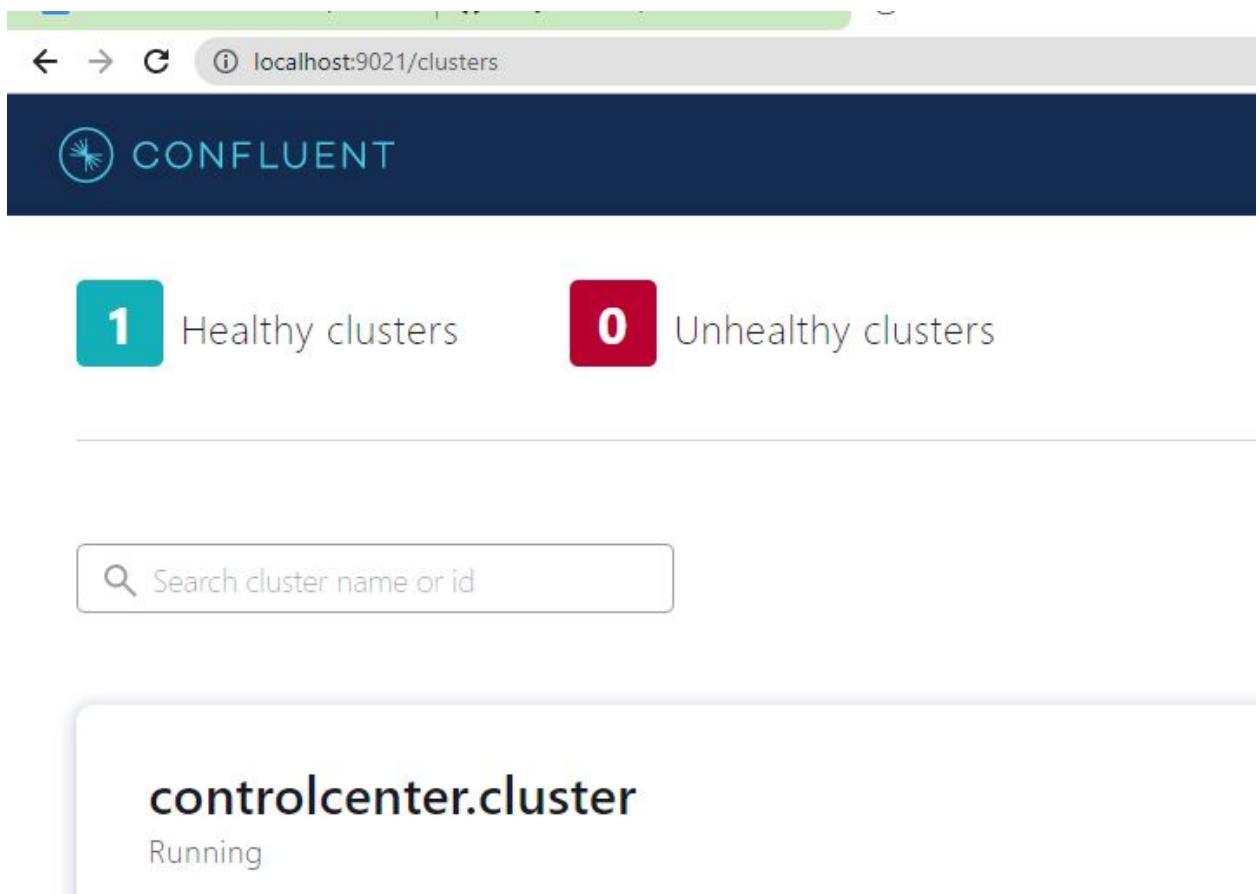
- You should see output as

- 

```
zookeeper is up-to-date
```

```
Starting broker ... done
schema-registry is up-to-date
rest-proxy is up-to-date
Starting connect ... done
Starting ksqldb-server ... done
Starting control-center ...
Starting ksql-datagen ...
Starting control-center ... done
```

- 
- Access confluent command center UI. VM is configured to do port forward of command center UI port. <http://localhost:9021/>
- 



Stopping Apache Kafka that is running in docker container

```
[root@kserver vagrant]# docker-compose stop
Stopping ksql-datagen ... done
Stopping control-center ... done
```

```
Stopping ksqldb-cli    ... done
Stopping ksqldb-server ... done
Stopping connect      ... done
Stopping rest-proxy   ... done
Stopping schema-registry ... done
Stopping broker       ... done
Stopping zookeeper    ... done
```

## Starting the container that was stopped

```
[root@kserver vagrant]# docker-compose start
Starting zookeeper    ... done
Starting broker       ... done
Starting schema-registry ... done
Starting connect      ... done
Starting ksqldb-server ... done
Starting control-center ... done
Starting ksqldb-cli   ... done
Starting ksql-datagen ... done
Starting rest-proxy   ... done
```

## Deleting Kafka environment

```
[root@kserver vagrant]# docker-compose down
Stopping control-center ... done
Stopping ksqldb-cli     ... done
Stopping ksqldb-server ... done
Stopping zookeeper      ... done
Removing ksql-datagen   ... done
Removing control-center ... done
Removing ksqldb-cli     ... done
Removing ksqldb-server ... done
Removing connect        ... done
Removing rest-proxy     ... done
Removing schema-registry ... done
Removing broker         ... done
Removing zookeeper      ... done
Removing network vagrant_default
```

## Troubleshooting

Checking state of kafka process that are running in docker container

```
[root@kserver vagrant]# docker-compose ps
```

Name	Command	State	Ports
broker	/etc/confluent/docker/run	Up	0.0.0.0:9092->9092/tcp, 0.0.0.0:9101->9101/tcp
connect	/etc/confluent/docker/run	Up	0.0.0.0:8083->8083/tcp, 9092/tcp
control-center	/etc/confluent/docker/run	Up	0.0.0.0:9021->9021/tcp
ksql-datagen	bash -c echo Waiting for K ...	Up	
ksqldb-cli	/bin/sh	Up	
ksqldb-server	/etc/confluent/docker/run	Up	0.0.0.0:8088->8088/tcp
rest-proxy	/etc/confluent/docker/run	Up	0.0.0.0:8082->8082/tcp
schema-registry	/etc/confluent/docker/run	Up	0.0.0.0:8081->8081/tcp
zookeeper	/etc/confluent/docker/run	Up	0.0.0.0:2181->2181/tcp, 2888/tcp, 3888/tcp

Stop the Linux virtual machine

```
CCD-Apache-Kafka-Certification-Examination-Notes> vagrant halt  
==> kserver: Attempting graceful shutdown of VM...
```

## Application Design

### Kafka's command line tools

Kafka provides several command-line interface (CLI) utilities that helps to manage topics, clusters etc

### Topic Operations

The **kafka-topics.sh** tool provides an easy way to perform most topic operations. This

allows you to create, modify, delete, and list information about topics in a cluster.

### Creating a New Topic in a cluster

```
docker-compose exec kafka-1 kafka-topics \  
  --create \  
  --bootstrap-server localhost:19092 \  
  --replication-factor 2 \  
  --partitions 8 \  
  --topic my-topic
```

Creates a new topic named my-topic with replication factor as 2 and partitions as 8

#### Key Points

- It is not recommend to have topic names starts with two underscore like \_\_my-topic
- Kafka internal topics starts with two underscores example \_\_**consumer\_offsets**
- In a single cluster we shouldn't have topic names with both underscores and period

### Increasing partitions of topic

```
# Increase partitions of a topic  
docker-compose exec kafka-1 kafka-topics \  
  --alter \  
  --bootstrap-server localhost:19092 \  
  --partitions 16 \  
  --topic my-topic
```

Increase the partitions of topic my-topic form 8 to 16

#### Key Points

- It is not best practice to increase the size of partitions once created.
- Topics that are having messages that is produced with Keys are difficult to change partitions

### Decreasing partitions of topic

It is not possible to reduce the number of partitions of a topic

### Deleting a Topic

```
# Delete a topic
docker-compose exec kafka-1 kafka-topics \
  --delete \
  --bootstrap-server localhost:19092 \
  --topic my-topic
```

Delete my-topic from the cluster

Key point

Delete a topic will also delete all the messages. This is a non reversible action.

Listing all topics in a Cluster

```
# list all topics in cluster
docker-compose exec kafka-1 kafka-topics \
  --list \
  --bootstrap-server localhost:19092
```

Describe all the topics

```
# Describe all the topics in cluster
docker-compose exec kafka-1 kafka-topics \
  --describe \
  --bootstrap-server localhost:19092
```

## Console Producer

The **kafka-console-producer.sh** tool can be used to write messages into a Kafka topic in your cluster

```
# Write messages to my-topic using console producer
docker-compose exec kafka-1 kafka-console-producer \
  --topic my-topic \
  --bootstrap-server localhost:19092 \
```

## Console Consumer

The **kafka-console-consumer.sh** tool is used to read messages from one or more topics



```
# Read messages from my-topic
docker-compose exec kafka-1 kafka-console-consumer \
  --topic my-topic \
  --bootstrap-server localhost:19092
```

By default the tool reads latest messages from topic

To read messages from beginning

```
# Read messages from my-topic from-beginning
docker-compose exec kafka-1 kafka-console-consumer \
  --topic my-topic \
  --from-beginning \
  --bootstrap-server localhost:19092
```

## Consumer Groups

The **kafka-consumer-groups.sh** tool can be used to list and describe both old and new consumer groups.

For old consumer the group information is maintained in Zookeeper

For new consumers the group information is maintained in Broker

List Consumer Groups in a cluster

```
# List new consumer groups
docker-compose exec kafka-1 kafka-consumer-groups \
  --list \
  --bootstrap-server localhost:19092
```

Describe a Consumer Group

```
# Describe a consumer group
docker-compose exec kafka-1 kafka-consumer-groups \
  --describe \
  --group console-consumer-51714 \
  --bootstrap-server localhost:19092
```

Development

Deployment/Testing/Monitoring

References

Confluent Docker	<a href="https://docs.confluent.io/platform/current/quickstart/cos-docker-quickstart.html">https://docs.confluent.io/platform/current/quickstart/cos-docker-quickstart.html</a>
Github url for confluent docker images	<a href="https://github.com/confluentinc/cp-all-in-one">https://github.com/confluentinc/cp-all-in-one</a>
Setting up multi node cluster	<a href="https://better-coding.com/building-apache-kafka-cluster-using-docker-compose-and-virtual-box/">https://better-coding.com/building-apache-kafka-cluster-using-docker-compose-and-virtual-box/</a>
Good Book	Kafka: The Definitive Guide- Real-Time Data and Stream Processing at Scale